

AirBnB Seattle Project

This project is part of the Udacity's Data Science Nanodegree program where I have chosen to write a blog post using AirBnB's Seattle Open Data.

The CRISP-DM (Cross Industry Process for Data Mining) process serves as a guide for the project and is as follows:

- Business Understanding
- Data Understanding
- Data Preparation
- Modeling
- Evaluation
- Deploy

Business Understanding

By exploring the AirBnB Seattle data, I hope to gain more understanding on the Seattle rental market, specifically, questions regarding:

- What property features determine the listing price?
- What property features determine its popularity? (with reviews per month as proxy)
- When is the most popular month to rent in Seattle?

Data Understanding

```
In [1]: # Import libraries
```

```

import glob
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

import helper_functions as helper

from datetime import datetime
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import make_scorer, fbeta_score, accuracy_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier

from sklearn.feature_extraction.text import CountVectorizer
from scipy.sparse import csr_matrix

pd.options.display.max_columns = 500
pd.options.display.max_rows = 500

```

```

In [2]: csvs = glob.glob('./input/seattle/*.csv')
        csvs

```

```

Out[2]: ['./input/seattle\\calendar.csv',
          './input/seattle\\listings.csv',
          './input/seattle\\reviews.csv']

```

```

In [3]: # Listings
        base = pd.read_csv(csvs[1])
        listings_df = base.copy()
        listings_df.head(2)

```

```

Out[3]:

```

id	listing_url	scrape_id	last_scraped	name	summary
----	-------------	-----------	--------------	------	---------

	id	listing_url	scrape_id	last_scraped	name	summary
0	241032	https://www.airbnb.com/rooms/241032	20160104002432	2016-01-04	Stylish Queen Anne Apartment	Na
1	953595	https://www.airbnb.com/rooms/953595	20160104002432	2016-01-04	Bright & Airy Queen Anne Apartment	Chemically sensitive? We've removed the irrita..

Features that are available for analysis: Including Non-Null count and Datatypes

In [4]: `listings_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3818 entries, 0 to 3817
Data columns (total 92 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    3818 non-null   int64
1   listing_url                          3818 non-null   object
2   scrape_id                            3818 non-null   int64
3   last_scraped                         3818 non-null   object
4   name                                 3818 non-null   object
5   summary                              3641 non-null   object
6   space                                3249 non-null   object
7   description                           3818 non-null   object
8   experiences_offered                   3818 non-null   object
9   neighborhood_overview                 2786 non-null   object
10  notes                                 2212 non-null   object
11  transit                               2884 non-null   object
```

12	thumbnail_url	3498	non-null	object
13	medium_url	3498	non-null	object
14	picture_url	3818	non-null	object
15	xl_picture_url	3498	non-null	object
16	host_id	3818	non-null	int64
17	host_url	3818	non-null	object
18	host_name	3816	non-null	object
19	host_since	3816	non-null	object
20	host_location	3810	non-null	object
21	host_about	2959	non-null	object
22	host_response_time	3295	non-null	object
23	host_response_rate	3295	non-null	object
24	host_acceptance_rate	3045	non-null	object
25	host_is_superhost	3816	non-null	object
26	host_thumbnail_url	3816	non-null	object
27	host_picture_url	3816	non-null	object
28	host_neighbourhood	3518	non-null	object
29	host_listings_count	3816	non-null	float64
30	host_total_listings_count	3816	non-null	float64
31	host_verifications	3818	non-null	object
32	host_has_profile_pic	3816	non-null	object
33	host_identity_verified	3816	non-null	object
34	street	3818	non-null	object
35	neighbourhood	3402	non-null	object
36	neighbourhood_cleansed	3818	non-null	object
37	neighbourhood_group_cleansed	3818	non-null	object
38	city	3818	non-null	object
39	state	3818	non-null	object
40	zipcode	3811	non-null	object
41	market	3818	non-null	object
42	smart_location	3818	non-null	object
43	country_code	3818	non-null	object
44	country	3818	non-null	object
45	latitude	3818	non-null	float64
46	longitude	3818	non-null	float64
47	is_location_exact	3818	non-null	object
48	property_type	3817	non-null	object
49	room_type	3818	non-null	object
50	accommodates	3818	non-null	int64

51	bathrooms	3802 non-null	float64
52	bedrooms	3812 non-null	float64
53	beds	3817 non-null	float64
54	bed_type	3818 non-null	object
55	amenities	3818 non-null	object
56	square_feet	97 non-null	float64
57	price	3818 non-null	object
58	weekly_price	2009 non-null	object
59	monthly_price	1517 non-null	object
60	security_deposit	1866 non-null	object
61	cleaning_fee	2788 non-null	object
62	guests_included	3818 non-null	int64
63	extra_people	3818 non-null	object
64	minimum_nights	3818 non-null	int64
65	maximum_nights	3818 non-null	int64
66	calendar_updated	3818 non-null	object
67	has_availability	3818 non-null	object
68	availability_30	3818 non-null	int64
69	availability_60	3818 non-null	int64
70	availability_90	3818 non-null	int64
71	availability_365	3818 non-null	int64
72	calendar_last_scraped	3818 non-null	object
73	number_of_reviews	3818 non-null	int64
74	first_review	3191 non-null	object
75	last_review	3191 non-null	object
76	review_scores_rating	3171 non-null	float64
77	review_scores_accuracy	3160 non-null	float64
78	review_scores_cleanliness	3165 non-null	float64
79	review_scores_checkin	3160 non-null	float64
80	review_scores_communication	3167 non-null	float64
81	review_scores_location	3163 non-null	float64
82	review_scores_value	3162 non-null	float64
83	requires_license	3818 non-null	object
84	license	0 non-null	float64
85	jurisdiction_names	3818 non-null	object
86	instant_bookable	3818 non-null	object
87	cancellation_policy	3818 non-null	object
88	require_guest_profile_picture	3818 non-null	object
89	require_guest_phone_verification	3818 non-null	object

```
90  calculated_host_listings_count    3818 non-null    int64
91  reviews_per_month                3191 non-null    float64
dtypes: float64(17), int64(13), object(62)
memory usage: 2.7+ MB
```

Data Preparation

A preliminary investigation reveals that further data cleansing and preprocessing is needed before it can be used as an input for our project pipeline.

The steps that needed to be undertaken are as per follows:

- Removing Unary and URL features except for 'thumbnail_url'.
- Removing features with many missing values : $\geq 21\%$ missing.
- Removing irrelevant features that lacks business justification or are too granular.
 - Text fields except for 'description'.
 - Fields with redundant information like 30,60,90 day availability, retain only availability_365.
 - Similarly with weekly, monthly price, retain only 'price' field.
 - The detailed split of review scores except the final rating.
 - More granular aspects of location.
- Remove special characters like '\$' and ',' from price features.
- Remove '%' character from percentage features.
- Convert thumbnail data as available or not-available.
- Convert host since from date to number of days.
- Convert amenities column to number of amenities.
- Impute Null as Zero for security deposit and cleaning fees.
- Count number of amenities from list.

Features that are Unary and URL

```
In [5]: unary_columns, url_columns = [], []
```

```

for i in listings_df.columns:
    if len(listings_df[i].unique())==1:
        print('Unary Column: ',i , listings_df[i].unique())
        unary_columns.append(i)
    if 'url' in i:
        url_columns = url_columns+[i]

```

```

Unary Column:  scrape_id [20160104002432]
Unary Column:  last_scraped ['2016-01-04']
Unary Column:  experiences_offered ['none']
Unary Column:  market ['Seattle']
Unary Column:  country_code ['US']
Unary Column:  country ['United States']
Unary Column:  has_availability ['t']
Unary Column:  calendar_last_scraped ['2016-01-04']
Unary Column:  requires_license ['f']
Unary Column:  license [nan]
Unary Column:  jurisdiction_names ['WASHINGTON']

```

In [6]: url_columns

```

Out[6]: ['listing_url',
         'thumbnail_url',
         'medium_url',
         'picture_url',
         'xl_picture_url',
         'host_url',
         'host_thumbnail_url',
         'host_picture_url']

```

In [7]: *# Drop all URL variables but convert [thumbnail_url] into binary*
listings_df[url_columns].isnull().sum()/len(listings_df)

```

Out[7]: listing_url          0.000000
         thumbnail_url      0.083814
         medium_url         0.083814
         picture_url        0.000000
         xl_picture_url     0.083814
         host_url           0.000000

```

```
host_thumbnail_url    0.000524
host_picture_url      0.000524
dtype: float64
```

```
In [8]: url_columns = list(set(url_columns) - {'thumbnail_url'})
```

Features that are Unary and URL Features with Many Missing Values

```
In [9]: percent_missing = listings_df.isnull().sum() * 100 / len(listings_df)
missing_value_df = pd.DataFrame({'column_name': listings_df.columns,
                                'percent_missing': round(percent_mi
ssing,2)})
missing_value_df.sort_values('percent_missing', inplace=True, ascending
=[False])
missing_value_df.reset_index(drop=True)
```

Out[9]:

	column_name	percent_missing
0	license	100.00
1	square_feet	97.46
2	monthly_price	60.27
3	security_deposit	51.13
4	weekly_price	47.38
5	notes	42.06
6	neighborhood_overview	27.03
7	cleaning_fee	26.98
8	transit	24.46
9	host_about	22.50
10	host_acceptance_rate	20.25
11	review_scores_accuracy	17.23
12	review_scores_checkin	17.23

13	review_scores_value	17.18
	column_name	percent_missing
14	review_scores_location	17.16
15	review_scores_cleanliness	17.10
16	review_scores_communication	17.05
17	review_scores_rating	16.95
18	last_review	16.42
19	first_review	16.42
20	reviews_per_month	16.42
21	space	14.90
22	host_response_rate	13.70
23	host_response_time	13.70
24	neighbourhood	10.90
25	thumbnail_url	8.38
26	medium_url	8.38
27	xl_picture_url	8.38
28	host_neighbourhood	7.86
29	summary	4.64
30	bathrooms	0.42
31	host_location	0.21
32	zipcode	0.18
33	bedrooms	0.16
34	host_name	0.05
35	host_listings_count	0.05
36	host_since	0.05
37	host_is_superhost	0.05

38	host_identity_verified	0.05
	column_name	percent_missing
39	host_picture_url	0.05
40	host_thumbnail_url	0.05
41	host_total_listings_count	0.05
42	host_has_profile_pic	0.05
43	property_type	0.03
44	beds	0.03
45	require_guest_profile_picture	0.00
46	calculated_host_listings_count	0.00
47	maximum_nights	0.00
48	calendar_updated	0.00
49	has_availability	0.00
50	require_guest_phone_verification	0.00
51	instant_bookable	0.00
52	availability_30	0.00
53	availability_60	0.00
54	availability_90	0.00
55	availability_365	0.00
56	calendar_last_scraped	0.00
57	number_of_reviews	0.00
58	cancellation_policy	0.00
59	jurisdiction_names	0.00
60	requires_license	0.00
61	extra_people	0.00
62	minimum_nights	0.00

63	id	0.00
	column_name	percent_missing
64	guests_included	0.00
65	price	0.00
66	scrape_id	0.00
67	last_scraped	0.00
68	name	0.00
69	description	0.00
70	experiences_offered	0.00
71	picture_url	0.00
72	host_id	0.00
73	host_url	0.00
74	host_verifications	0.00
75	street	0.00
76	neighbourhood_cleansed	0.00
77	neighbourhood_group_cleansed	0.00
78	city	0.00
79	state	0.00
80	market	0.00
81	smart_location	0.00
82	country_code	0.00
83	country	0.00
84	latitude	0.00
85	listing_url	0.00
86	is_location_exact	0.00
87	room_type	0.00

88	accommodates	0.00
column_name		percent_missing
89	bed_type	0.00
90	amenities	0.00
91	longitude	0.00

```
In [10]: missing_value_df = missing_value_df[missing_value_df['percent_missing']
        >=21.00].reset_index(drop=True)
        missing_value_df
```

Out[10]:

column_name		percent_missing
0	license	100.00
1	square_feet	97.46
2	monthly_price	60.27
3	security_deposit	51.13
4	weekly_price	47.38
5	notes	42.06
6	neighborhood_overview	27.03
7	cleaning_fee	26.98
8	transit	24.46
9	host_about	22.50

Impute NULL to be ZERO for security deposit and cleaning fee

This is based on the assumption that hosts do not input any value when they are not charging the fees

```
In [11]: # Assume NULL to be ZERO in security_deposit, cleaning_fee
```

```
missing_columns = missing_value_df['column_name'].values.tolist()
missing_columns = list(set(missing_columns) - {'security_deposit', 'cleaning_fee'})
missing_columns
```

```
Out[11]: ['license',
          'weekly_price',
          'monthly_price',
          'notes',
          'square_feet',
          'host_about',
          'transit',
          'neighborhood_overview']
```

Removing Features that are obviously irrelevant (as described above)

```
In [12]: irrelevant_columns = [
        , 'id'
        , 'name'
        , 'summary'
        , 'space'
        , 'host_id'
        , 'host_name'
        , 'host_location'
        , 'host_neighbourhood'
        , 'host_listings_count'
        , 'host_total_listings_count'
        , 'host_verifications'
        , 'street'
        , 'neighbourhood'
        , 'neighbourhood_cleansed'
        , 'city'
        , 'state'
        , 'zipcode'
        , 'smart_location'
        , 'latitude'
        , 'longitude'
        , 'is_location_exact'
        , 'minimum_nights'
```

```
, 'maximum_nights'  
, 'calendar_updated'  
, 'availability_30'  
, 'availability_60'  
, 'availability_90'  
, 'first_review'  
, 'last_review'  
, 'require_guest_profile_picture'  
, 'require_guest_phone_verification'  
, 'calculated_host_listings_count']
```

Features that are Unary and URL Consolidate List of Features that are to be removed

```
In [13]: remove_columns = unary_columns + url_columns + missing_columns + irrelevant_columns  
remove_columns
```

```
Out[13]: ['scrape_id',  
'last_scraped',  
'experiences_offered',  
'market',  
'country_code',  
'country',  
'has_availability',  
'calendar_last_scraped',  
'requires_license',  
'license',  
'jurisdiction_names',  
'xl_picture_url',  
'listing_url',  
'host_url',  
'picture_url',  
'medium_url',  
'host_thumbnail_url',  
'host_picture_url',  
'license',  
'weekly_price',  
'monthly_price']
```

```
monthly_price ,
'notes',
'square_feet',
'host_about',
'transit',
'neighborhood_overview',
'id',
'name',
'summary',
'space',
'host_id',
'host_name',
'host_location',
'host_neighbourhood',
'host_listings_count',
'host_total_listings_count',
'host_verifications',
'street',
'neighbourhood',
'neighbourhood_cleansed',
'city',
'state',
'zipcode',
'smart_location',
'latitude',
'longitude',
'is_location_exact',
'minimum_nights',
'maximum_nights',
'calendar_updated',
'availability_30',
'availability_60',
'availability_90',
'first_review',
'last_review',
'require_guest_profile_picture',
'require_guest_phone_verification',
'calculated_host_listings_count']
```

```
In [14]: listings_df.drop(listings_df[remove_columns], axis=1, inplace=True)
```

```
In [15]: listings_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3818 entries, 0 to 3817
```

```
Data columns (total 35 columns):
```

#	Column	Non-Null Count	Dtype
0	description	3818 non-null	object
1	thumbnail_url	3498 non-null	object
2	host_since	3816 non-null	object
3	host_response_time	3295 non-null	object
4	host_response_rate	3295 non-null	object
5	host_acceptance_rate	3045 non-null	object
6	host_is_superhost	3816 non-null	object
7	host_has_profile_pic	3816 non-null	object
8	host_identity_verified	3816 non-null	object
9	neighbourhood_group_cleansed	3818 non-null	object
10	property_type	3817 non-null	object
11	room_type	3818 non-null	object
12	accommodates	3818 non-null	int64
13	bathrooms	3802 non-null	float64
14	bedrooms	3812 non-null	float64
15	beds	3817 non-null	float64
16	bed_type	3818 non-null	object
17	amenities	3818 non-null	object
18	price	3818 non-null	object
19	security_deposit	1866 non-null	object
20	cleaning_fee	2788 non-null	object
21	guests_included	3818 non-null	int64
22	extra_people	3818 non-null	object
23	availability_365	3818 non-null	int64
24	number_of_reviews	3818 non-null	int64
25	review_scores_rating	3171 non-null	float64
26	review_scores_accuracy	3160 non-null	float64
27	review_scores_cleanliness	3165 non-null	float64
28	review_scores_checkin	3160 non-null	float64
29	review_scores_communication	3167 non-null	float64


```
30 review_scores_location      3163 non-null    float64
31 review_scores_value         3162 non-null    float64
32 instant_bookable            3818 non-null    object
33 cancellation_policy          3818 non-null    object
34 reviews_per_month           3191 non-null    float64
dtypes: float64(11), int64(4), object(20)
memory usage: 1.0+ MB
```

Invoke Helper functions to further clean the dataframe

```
In [16]: seattle_df = helper.clean_data(listings_df)
```

Data Preprocessing: Price

Based on the questions above two different set of features were required to carry out the analysis.

- Keep remaining features except property text description for predicting price.
- Use property text description alone for price prediction.

```
In [17]: seattle_df.drop('description', axis=1, inplace=True)
seattle_df.head()
```

Out[17]:

	host_since	host_response_time	host_response_rate	host_acceptance_rate	host_is_superhost
0	1607.0	within a few hours	96.0	100.0	0
1	1047.0	within an hour	98.0	100.0	1
2	571.0	within a few hours	67.0	100.0	0
3	789.0	NaN	NaN	NaN	0

	host_since	host_response_time	host_response_rate	host_acceptance_rate	host_is_superhost
4	1497.0	within an hour	100.0	NaN	0

In [18]: `seattle_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3818 entries, 0 to 3817
Data columns (total 34 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   host_since                               3816 non-null   float64
1   host_response_time                       3295 non-null   object
2   host_response_rate                       3295 non-null   float64
3   host_acceptance_rate                     3045 non-null   float64
4   host_is_superhost                        3818 non-null   int64
5   host_has_profile_pic                     3818 non-null   int64
6   host_identity_verified                   3818 non-null   int64
7   neighbourhood_group_cleansed             3818 non-null   object
8   property_type                            3817 non-null   object
9   room_type                               3818 non-null   object
10  accommodates                             3818 non-null   int64
11  bathrooms                               3802 non-null   float64
12  bedrooms                               3812 non-null   float64
13  beds                                   3817 non-null   float64
14  bed_type                                3818 non-null   object
15  price                                   3818 non-null   float64
16  security_deposit                         3818 non-null   float64
17  cleaning_fee                             3818 non-null   float64
18  guests_included                         3818 non-null   int64
19  extra_people                             3818 non-null   float64
20  availability_365                         3818 non-null   int64
21  number_of_reviews                       3818 non-null   int64
22  review_scores_rating                     3171 non-null   float64
23  review_scores_accuracy                   3160 non-null   float64
24  review_scores_cleanliness                3165 non-null   float64
25  review_scores_checkin                    3160 non-null   float64
```

```

26 review_scores_communication 3167 non-null float64
27 review_scores_location      3163 non-null float64
28 review_scores_value         3162 non-null float64
29 instant_bookable            3818 non-null int64
30 cancellation_policy          3818 non-null int64
31 reviews_per_month           3191 non-null float64
32 has_thumbnail_url           3818 non-null int32
33 total_amenities              3818 non-null int64
dtypes: float64(18), int32(1), int64(10), object(5)
memory usage: 999.4+ KB

```

Splitting Data into Features/Label and Dividing Label/Price feature into (High,Low)=(1,0)

```

In [19]: label_df = np.where(seattle_df['price'] > seattle_df['price'].median(
(), 0, 1)
feature_df = seattle_df.drop(['price'], axis=1)

```

Impute Missing Categorical Values with 'Most Frequent' and Apply Min-Max Feature Scaling

```

In [20]: scaled_df = helper.process_features(feature_df)

```

```

In [21]: scaled_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3818 entries, 0 to 3817
Data columns (total 73 columns):
#   Column                                     Non-Null Count
Dtype
---  -----
0   host_since                                3818 non-null
float64
1   host_response_rate                       3818 non-null
float64

```

2	host_acceptance_rate	3818 non-null
float64		
3	host_is_superhost	3818 non-null
float64		
4	host_has_profile_pic	3818 non-null
float64		
5	host_identity_verified	3818 non-null
float64		
6	accommodates	3818 non-null
float64		
7	bathrooms	3818 non-null
float64		
8	bedrooms	3818 non-null
float64		
9	beds	3818 non-null
float64		
10	security_deposit	3818 non-null
float64		
11	cleaning_fee	3818 non-null
float64		
12	guests_included	3818 non-null
float64		
13	extra_people	3818 non-null
float64		
14	availability_365	3818 non-null
float64		
15	number_of_reviews	3818 non-null
float64		
16	review_scores_rating	3818 non-null
float64		
17	review_scores_accuracy	3818 non-null
float64		
18	review_scores_cleanliness	3818 non-null
float64		
19	review_scores_checkin	3818 non-null
float64		
20	review_scores_communication	3818 non-null
float64		
21	review_scores_location	3818 non-null

float64		
22	review_scores_value	3818 non-null
float64		
23	instant_bookable	3818 non-null
float64		
24	cancellation_policy	3818 non-null
float64		
25	reviews_per_month	3818 non-null
float64		
26	has_thumbnail_url	3818 non-null
float64		
27	total_amenities	3818 non-null
float64		
28	host_response_time_a_few_days_or_more	3818 non-null
float64		
29	host_response_time_within_a_day	3818 non-null
float64		
30	host_response_time_within_a_few_hours	3818 non-null
float64		
31	host_response_time_within_an_hour	3818 non-null
float64		
32	neighbourhood_group_cleansed_Ballard	3818 non-null
float64		
33	neighbourhood_group_cleansed_Beacon_Hill	3818 non-null
float64		
34	neighbourhood_group_cleansed_Capitol_Hill	3818 non-null
float64		
35	neighbourhood_group_cleansed_Cascade	3818 non-null
float64		
36	neighbourhood_group_cleansed_Central_Area	3818 non-null
float64		
37	neighbourhood_group_cleansed_Delridge	3818 non-null
float64		
38	neighbourhood_group_cleansed_Downtown	3818 non-null
float64		
39	neighbourhood_group_cleansed_Interbay	3818 non-null
float64		
40	neighbourhood_group_cleansed_Lake_City	3818 non-null
float64		

41	neighbourhood_group_cleansed_Magnolia	3818	non-null
	float64		
42	neighbourhood_group_cleansed_Northgate	3818	non-null
	float64		
43	neighbourhood_group_cleansed_Other_neighborhoods	3818	non-null
	float64		
44	neighbourhood_group_cleansed_Queen_Anne	3818	non-null
	float64		
45	neighbourhood_group_cleansed_Rainier_Valley	3818	non-null
	float64		
46	neighbourhood_group_cleansed_Seward_Park	3818	non-null
	float64		
47	neighbourhood_group_cleansed_University_District	3818	non-null
	float64		
48	neighbourhood_group_cleansed_West_Seattle	3818	non-null
	float64		
49	property_type_Apartment	3818	non-null
	float64		
50	property_type_Bed_and_Breakfast	3818	non-null
	float64		
51	property_type_Boat	3818	non-null
	float64		
52	property_type_Bungalow	3818	non-null
	float64		
53	property_type_Cabin	3818	non-null
	float64		
54	property_type_Camper_RV	3818	non-null
	float64		
55	property_type_Chalet	3818	non-null
	float64		
56	property_type_Condominium	3818	non-null
	float64		
57	property_type_Dorm	3818	non-null
	float64		
58	property_type_House	3818	non-null
	float64		
59	property_type_Loft	3818	non-null
	float64		
60	property_type_Other	3818	non-null

```
float64
  61  property_type_Tent                3818 non-null
float64
  62  property_type_Townhouse           3818 non-null
float64
  63  property_type_Treehouse           3818 non-null
float64
  64  property_type_Yurt                 3818 non-null
float64
  65  room_type_Entire_home_apt          3818 non-null
float64
  66  room_type_Private_room             3818 non-null
float64
  67  room_type_Shared_room              3818 non-null
float64
  68  bed_type_Airbed                    3818 non-null
float64
  69  bed_type_Couch                     3818 non-null
float64
  70  bed_type_Futon                     3818 non-null
float64
  71  bed_type_Pull_out_Sofa             3818 non-null
float64
  72  bed_type_Real_Bed                  3818 non-null
float64
dtypes: float64(73)
memory usage: 2.1 MB
```

Modeling and Evaluation: Price

What property features determine the listing price?

AdaBoostClassifier with decision tree base estimator and optimized with GridSearchCV.

```
In [22]: # Initialize the classifier
         clf = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), random_state=42)
```

```
m_state=123)

# Create the parameters list to tune, using a dictionary.
parameters = {"n_estimators": [10, 50, 100],
               "learning_rate": [0.005, .01, 0.05, 0.1],
               'base_estimator__min_samples_split' : [2, 4, 6, 8],
               'base_estimator__max_depth' : [2, 4, 6, 8]}
```

In [23]: `best_clf, X_train, X_test, y_train, y_test = helper.boost_classifier(clf, parameters, scaled_df, label_df)`

```
C:\Users\tai_j\Anaconda3\lib\site-packages\sklearn\utils\validation.py:
71: FutureWarning: Pass scoring=make_scorer(fbeta_score, beta=0.5) as k
keyword args. From version 0.25 passing these as positional arguments wi
ll result in an error
FutureWarning)
```

In [24]: `print(best_clf)`

```
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=2,
                                                         min_samples_sp
lit=6),
                  learning_rate=0.1, n_estimators=100, random_state=12
3)
```

In [25]: `# Unoptimized model`
`test_accuracy, train_accuracy = helper.prediction_scores(clf, X_train,`
`X_test, y_train, y_test)`
`helper.print_scores(test_accuracy, train_accuracy)`

```
Accuracy score on testing data: 0.7304
Accuracy score on training data: 1.0000
```

In [26]: `# Optimized model`
`test_accuracy, train_accuracy = helper.prediction_scores(best_clf, X_tr`
`ain, X_test, y_train, y_test)`
`helper.print_scores(test_accuracy, train_accuracy)`

Accuracy score on testing data: 0.8272
Accuracy score on training data: 0.8333

Top 10 important features

```
In [27]: # Extract the feature importances using .feature_importances_
importances = best_clf.feature_importances_
indices = np.argsort(importances[::-1])
print('Top 10 Important Features')
display(X_train.columns.values[indices[:10]])
```

Top 10 Important Features

```
array(['room_type_Entire_home_aprt', 'accommodates', 'availability_365',
      'reviews_per_month', 'neighbourhood_group_cleansed_Downtown',
      'cleaning_fee', 'bedrooms', 'property_type_Bed_and_Breakfast',
      'neighbourhood_group_cleansed_Capitol_Hill', 'host_since'],
      dtype=object)
```

Here, we can see attributes relating to its size/type and location are relatively more important than reviews per month and experience of host, which makes sense. The bigger the unit, the more services/amenities provided, the more the host is going to charge.

Modeling and Evaluation: Popularity

What property features determine its popularity?

- Using number of reviews per month as a proxy for popularity.
- Similar data preparation as for previous price modeling.

```
In [28]: # Remove records with Null values in 'reviews_per_month'
revpm_df = seattle_df[seattle_df['reviews_per_month'].isnull()
!= True]
```

```
# Split Data into features and labels
label_revpm_df = np.where(revpm_df['reviews_per_month'] > revpm_df['reviews_per_month'].median(), 0, 1)
feature_revpm_df = revpm_df.drop(['reviews_per_month'], axis=1)

# Feature Scaling
scaled_revpm_df = helper.process_features(feature_revpm_df)
scaled_revpm_df.head()
```

Out[28]:

	host_since	host_response_rate	host_acceptance_rate	host_is_superhost	host_has_profile_pic
0	0.611305	0.951807	1.0	0.0	1.0
1	0.394503	0.975904	1.0	1.0	1.0
2	0.210221	0.602410	1.0	0.0	1.0
4	0.568719	1.000000	1.0	0.0	1.0
5	0.699961	1.000000	1.0	0.0	1.0

AdaBoostClassifier with decision tree base estimator. The model was optimized using GridSearchCV.

```
In [29]: # Initialize the classifier
clf = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), random_state=123)

# Create the parameters list to tune, using a dictionary.
parameters = {"n_estimators": [10, 50, 100],
              "learning_rate": [0.005, .01, 0.05, 0.1],
              'base_estimator__min_samples_split' : [2, 4, 6, 8],
              'base_estimator__max_depth' : [2, 4, 6, 8]}
```

```
In [30]: best_clf, X_train, X_test, y_train, y_test = helper.boost_classifier(clf, parameters, scaled_revpm_df, label_revpm_df)
```

```
C:\Users\tai_j\Anaconda3\lib\site-packages\sklearn\utils\validation.py:
71: FutureWarning: Pass scoring=make_scorer(fbeta_score, beta=0.5) as k
keyword args. From version 0.25 passing these as positional arguments wi
ll result in an error
FutureWarning)
```

In [31]: `print(best_clf)`

```
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=2,
                                                         min_samples_sp
lit=8),
                  learning_rate=0.1, n_estimators=100, random_state=12
3)
```

In [32]: `# Unoptimized model`
`test_accuracy, train_accuracy = helper.prediction_scores(clf, X_train,`
`X_test, y_train, y_test)`
`helper.print_scores(test_accuracy, train_accuracy)`

```
Accuracy score on testing data: 0.7825
Accuracy score on training data: 1.0000
```

In [33]: `# Optimized model`
`test_accuracy, train_accuracy = helper.prediction_scores(best_clf, X_tr`
`ain, X_test, y_train, y_test)`
`helper.print_scores(test_accuracy, train_accuracy)`

```
Accuracy score on testing data: 0.8607
Accuracy score on training data: 0.8687
```

In [34]: `# Extract the feature importances using .feature_importances_`
`importances = best_clf.feature_importances_`
`indices = np.argsort(importances)[::-1]`
`print('Top 10 Important Features')`
`display(X_train.columns.values[indices[:10]])`

```
Top 10 Important Features
```

```
array(['number_of_reviews', 'host_since', 'availability_365',
```

```
'cleaning_fee', 'price', 'total_amenities',  
'host_response_time_within_an_hour', 'bedrooms',  
'instant_bookable', 'host_is_superhost'], dtype=object)
```

Here, we can see number of reviews, experience of host/superhost, property availability are more relatively more dominant relative to attributes of the properties itself, unlike what we saw previously. This seems to lend weight to the assumption that popularity is in part driven by virality of the property.

When is the most popular month to rent in Seattle?

- Jul, Aug and Sep are the best period to maximise revenue. Before May are the best time for maintenance work Oct to Dec is a good time to take a break and enjoy the holidays.

```
In [35]: # Listings Dataframe  
base = pd.read_csv(csvs[1])  
listings_df = base.copy()  
listings_df.rename(columns={'id': 'listing_id'}, inplace=True)  
  
# Reviews Dataframe  
base = pd.read_csv(csvs[2])  
reviews_df = base.copy()  
reviews_df = reviews_df[['id', 'listing_id', 'date']]  
  
# Datetime Conversion  
reviews_df['date'] = pd.to_datetime(reviews_df['date'])  
reviews_df.head()
```

Out[35]:

	id	listing_id	date
0	38917982	7202016	2015-07-19
1	39087409	7202016	2015-07-20

	id	listing_id	date
2	39820030	7202016	2015-07-26
3	40813543	7202016	2015-08-02
4	41986501	7202016	2015-08-10

```
In [45]: # Bookings Dataframe
bookings_df = pd.merge(reviews_df, listings_df, on='listing_id')
bookings_df['estimated_revenue'] = bookings_df['price'] * bookings_df['minimum_nights']
bookings_df.head(20)
```

Out[45]:

	id	listing_id	date	minimum_nights	price	estimated_revenue
0	38917982	7202016	2015-07-19	2	75.0	150.0
1	39087409	7202016	2015-07-20	2	75.0	150.0
2	39820030	7202016	2015-07-26	2	75.0	150.0
3	40813543	7202016	2015-08-02	2	75.0	150.0
4	41986501	7202016	2015-08-10	2	75.0	150.0
5	43979139	7202016	2015-08-23	2	75.0	150.0
6	45265631	7202016	2015-09-01	2	75.0	150.0
7	46749120	7202016	2015-09-13	2	75.0	150.0
8	47783346	7202016	2015-09-21	2	75.0	150.0
9	48388999	7202016	2015-09-26	2	75.0	150.0
10	49441269	7202016	2015-10-04	2	75.0	150.0
11	50490194	7202016	2015-10-12	2	75.0	150.0
12	53862449	7202016	2015-11-13	2	75.0	150.0
13	54562283	7202016	2015-11-21	2	75.0	150.0
14	55212826	7202016	2015-11-29	2	75.0	150.0
15	58268184	7202016	2016-01-02	2	75.0	150.0

	id	listing_id	date	minimum_nights	price	estimated_revenue
16	20798623	3946674	2014-10-05	1	90.0	90.0
17	21224862	3946674	2014-10-13	1	90.0	90.0
18	22877803	3946674	2014-11-16	1	90.0	90.0
19	22938377	3946674	2014-11-17	1	90.0	90.0

```
In [49]: # get revenue by listings
listings_df=listings_df[['listing_id','minimum_nights','price']]
listings_df['price'] = listings_df['price'].map(lambda x: helper.convert_to_price(x))

listings_df_revenue = bookings_df[['listing_id','estimated_revenue']].groupby(['listing_id']).sum()
listings_df = pd.merge(listings_df, listings_df_revenue, on='listing_id', how='left')
listings_df.at[listings_df['estimated_revenue'].isnull(), 'estimated_revenue'] = 0
listings_df
```

C:\Users\tai_j\Anaconda3\lib\site-packages\ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
This is separate from the ipykernel package so we can avoid doing imports until

Out[49]:

	listing_id	minimum_nights	price	estimated_revenue
0	241032	1	85.0	17595.0
1	953595	2	150.0	12900.0
2	3308979	4	975.0	78000.0

	listing_id	minimum_nights	price	estimated_revenue
3	7421966	1	100.0	0.0
4	278830	1	450.0	17100.0
...
3813	8101950	3	359.0	1077.0
3814	8902327	2	79.0	316.0
3815	10267360	1	93.0	0.0
3816	9604740	3	99.0	0.0
3817	10208623	1	87.0	0.0

3818 rows × 4 columns

```
In [50]: plt.figure(figsize=(15, 5))

# # bookings by month
plotdata1 = reviews_df[['date']].groupby(reviews_df["date"].dt.month).count()
plotdata1.rename(columns={'date': '# of bookings'}, inplace=True)

ax = plt.subplot(1, 3, 1)
ax.set_title("# bookings by month")
plt.bar(plotdata1.index, plotdata1['# of bookings'])

# revenue by month
plotdata2 = bookings_df[['date', 'estimated_revenue']].groupby(bookings_df["date"].dt.month).sum()
plotdata2.rename(columns={'estimated_revenue': 'revenue'}, inplace=True)

ax = plt.subplot(1, 3, 2)
ax.set_title("revenue by month")
plt.bar(plotdata2.index, plotdata2['revenue'])

# avg booking price by month
plotdata3 = pd.concat([plotdata1, plotdata2], axis=1)
```


```

plotdata3['avg booking price'] = plotdata3['revenue'] / plotdata3['# of
bookings']
plotdata3.head()

ax = plt.subplot(1, 3, 3)
ax.set_title("avg booking price by month")
plt.bar(plotdata3.index, plotdata3['avg booking price'])

_ = plt.plot()

```



The impact of seasonality is unmistakable: August is the most popular month of the year by far compared to other months. The lower average booking price is more than made up by the volume of bookings. This lends credence to the idea that we should be mindful of attributes not captured by our data, but one that is significant enough to appear in a simple plot like this.

Conclusion

This simple and straightforward data-driven investigation of the Seattle Airbnb market has provided us with insights that may be helpful to anything keen to be involved in the property rental business.

A data-driven strategy would be important for any potential Airbnb host to get the important things right, and to fine-tune their existing listings to obtain higher premiums for their listing.

prices.

As per our investigation, there are factors that Airbnb hosts can tune: like (1) setting optimal prices/cleaning fees, (2) quick to reply, and (3) making the lodging available for more days of the year, and to some extent: (4) more amenities (perhaps adding on wifi, free breakfast, etc).

And there are factors that will be more relevant to people who are investing in a property or who have the resources to modify their lodgings, like: (1) making sure the property is situated in Downtown/Capitol Hill, (2) more rooms, (3) making whole apartments available instead of single-rooms, (4) enlarging the unit to accommodate more.

Understanding all these factors in depth will definitely provide an edge to anyone competing in the property rental market.